Abram Yorde

Englewood, OH | abram.yorde@outlook.com | (937) 507-1958 | Abram Yorde | AbramYorde

Objective

Enable internal and external customer success via data and analytics solutions and models.

Ensure these customer's success and guide them from reactionary to proactive and guided decision making.

Education

MS **Georgia Institute of Technology**, Analytics — Atlanta, GA
- GPA: 3.66 / 4.00

Jan 2019 – Dec 2021

BS **Wright State University**, Mechanical Engineering — Dayton, OH
- GPA: 3.71 / 4.00

Jan 2015 – Dec 2017

Experience

Copeland, Senior Data Engineer — Remote
- Lorem Epsum

Apr 2024 – present
1 year 9 months

Copeland, Senior Data Scientist — Sidney, OH
- Lorem Epsum

Jan 2023 – Aug 2024
1 year 8 months

Emerson Climate Technologies, Data Scientist — Sidney, OH
- Lorem Epsum

Jan 2020 – Jan 2023
3 years 1 month

Emerson Climate Technologies, Refrigeration Project Engineer - Data Science — Sidney, OH
- Created on-device neural network compression pipeline deployed across 50M+ devices
- Filed 2 patents on efficient model quantization techniques for edge inference

Jan 2018 – Jan 2020
2 years 1 month

Emerson Climate Technologies, Refrigeration Internships (Multiple Departments) — Sidney, OH
- Implemented novel self-supervised learning framework for low-resource language modeling
- Research integrated into Azure Cognitive Services, reducing training data requirements by 60%

Aug 2015 – Jan 2018
2 years 6 months

Projects

FlashInfer — Jan 2023 – present

Open-source library for high-performance LLM inference kernels
- Achieved 2.8x speedup over baseline attention implementations on A100 GPUs
- Adopted by 3 major AI labs, 8,500+ GitHub stars, 200+ contributors

NeuralPrune — Jan 2021

Automated neural network pruning toolkit with differentiable masks
- Reduced model size by 90% with less than 1% accuracy degradation on ImageNet
- Featured in PyTorch ecosystem tools, 4,200+ GitHub stars

Selected Honors

- MIT Technology Review 35 Under 35 Innovators (2024)
- Forbes 30 Under 30 in Enterprise Technology (2024)
- ACM Doctoral Dissertation Award Honorable Mention (2023)
- Google PhD Fellowship in Machine Learning (2020 – 2023)
- Fulbright Scholarship for Graduate Studies (2018)

Skills

Languages: Python, C++, CUDA, Rust, Julia

ML Frameworks: PyTorch, JAX, TensorFlow, Triton, ONNX

Infrastructure: Kubernetes, Ray, distributed training, AWS, GCP

Research Areas: Neural architecture search, model compression, efficient inference, multi-agent RL

Patents

1. Adaptive Quantization for Neural Network Inference on Edge Devices (US Patent 11,234,567)
2. Dynamic Sparsity Patterns for Efficient Transformer Attention (US Patent 11,345,678)
3. Hardware-Aware Neural Architecture Search Method (US Patent 11,456,789)

Invited Talks

4. Scaling Laws for Efficient Inference — Stanford HAI Symposium (2024)
3. Building AI Infrastructure for the Next Decade — TechCrunch Disrupt (2024)
2. From Research to Production: Lessons in ML Systems — NeurIPS Workshop (2023)
1. Efficient Deep Learning: A Practitioner's Perspective — Google Tech Talk (2022)

Any Section Title

You can use any section title you want.

You can choose any entry type for the section: `TextEntry`, `ExperienceEntry`, `EducationEntry`, `PublicationEntry`, `BulletEntry`, `NumberedEntry`, or `ReversedNumberedEntry`.

Markdown syntax is supported everywhere.

The `design` field in YAML gives you control over almost any aspect of your CV design.

See the documentation for more details.